

# Any-Precision LLM: Low-Cost Deployment of Multiple, Different-Sized LLMs

**Yeonhong Park**

✉ [ilil96@snu.ac.kr](mailto:ilil96@snu.ac.kr)

**Jake Hyun**

✉ [jakehyun@snu.ac.kr](mailto:jakehyun@snu.ac.kr)

**SangLyul Cho**

✉ [chosanglyul@snu.ac.kr](mailto:chosanglyul@snu.ac.kr)

**Bonggeun Sim**


✉ [bg.sim@snu.ac.kr](mailto:bg.sim@snu.ac.kr)

**Jae W. Lee**

✉ [jaewlee@snu.ac.kr](mailto:jaewlee@snu.ac.kr)



# Often, We Need Multiple LLMs of Different Sizes




# Often, We Need Multiple LLMs of Different Sizes

**Scenario #1:** Different queries with different latency requirements




# Often, We Need Multiple LLMs of Different Sizes

## Scenario #1: Different queries with different latency requirements




# Often, We Need Multiple LLMs of Different Sizes

**Scenario #2:** Dynamic fluctuation of system load and power budget



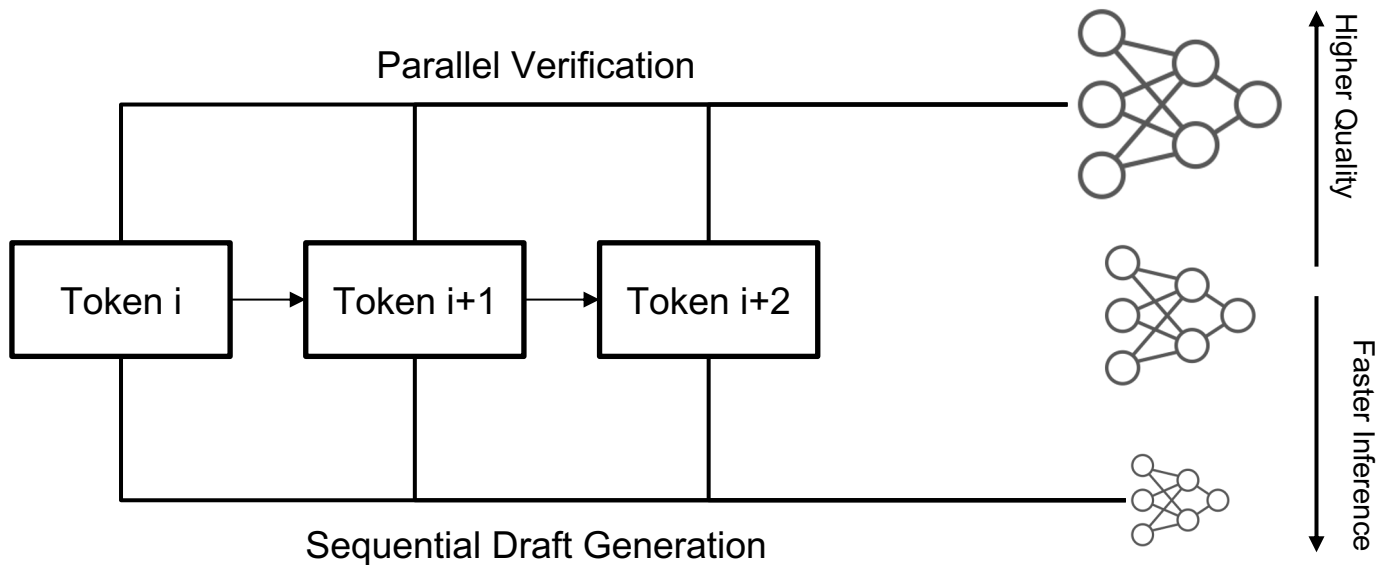
# Often, We Need Multiple LLMs of Different Sizes

**Scenario #2:** Dynamic fluctuation of system load and power budget



# Often, We Need Multiple LLMs of Different Sizes

## Scenario #3: Speculative decoding




# Research Question

*How can we deploy multiple, different-sized LLMs in a memory-efficient way?*






# Solution: Any-Precision Quantization (Yu et al., 2021<sup>[1]</sup>)

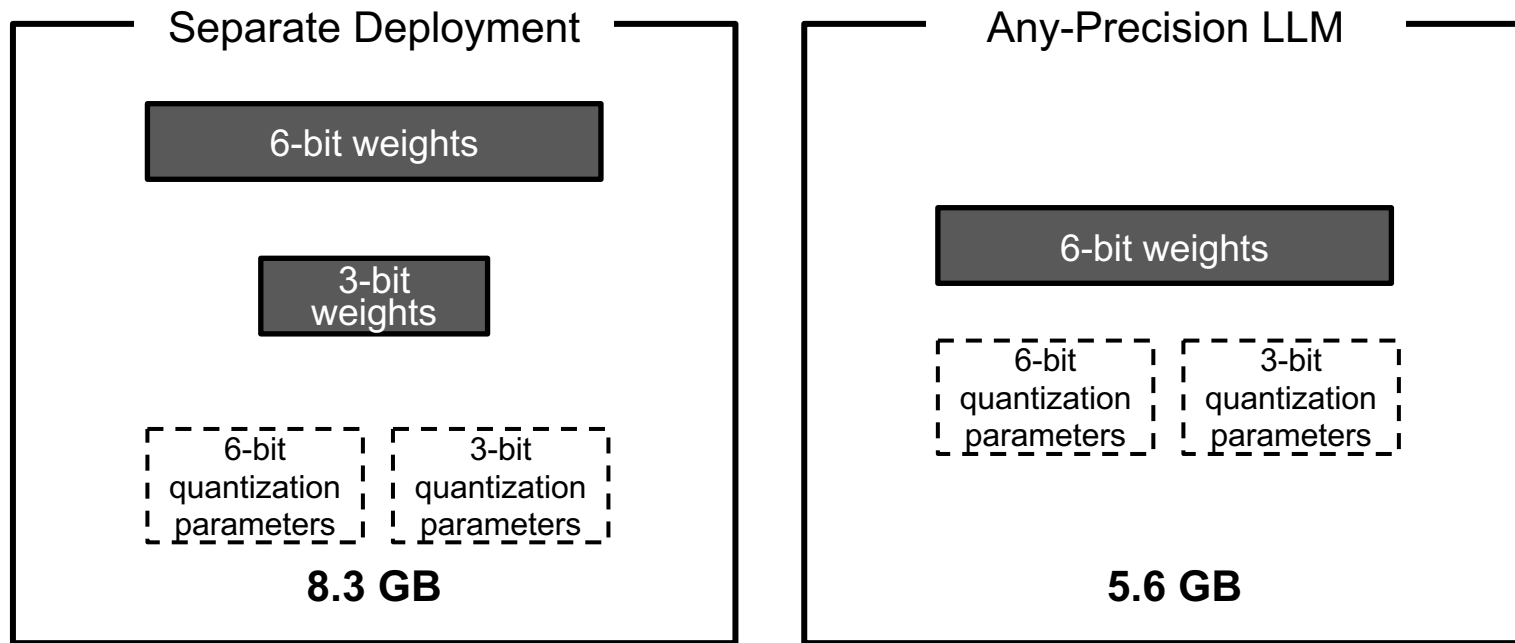


# Solution: Any-Precision Quantization (Yu et al., 2021<sup>[1]</sup>)



# Potential Memory Saving of Any-Precision LLM (Llama-2-7B)

Supported bit-widths: {3, 6}

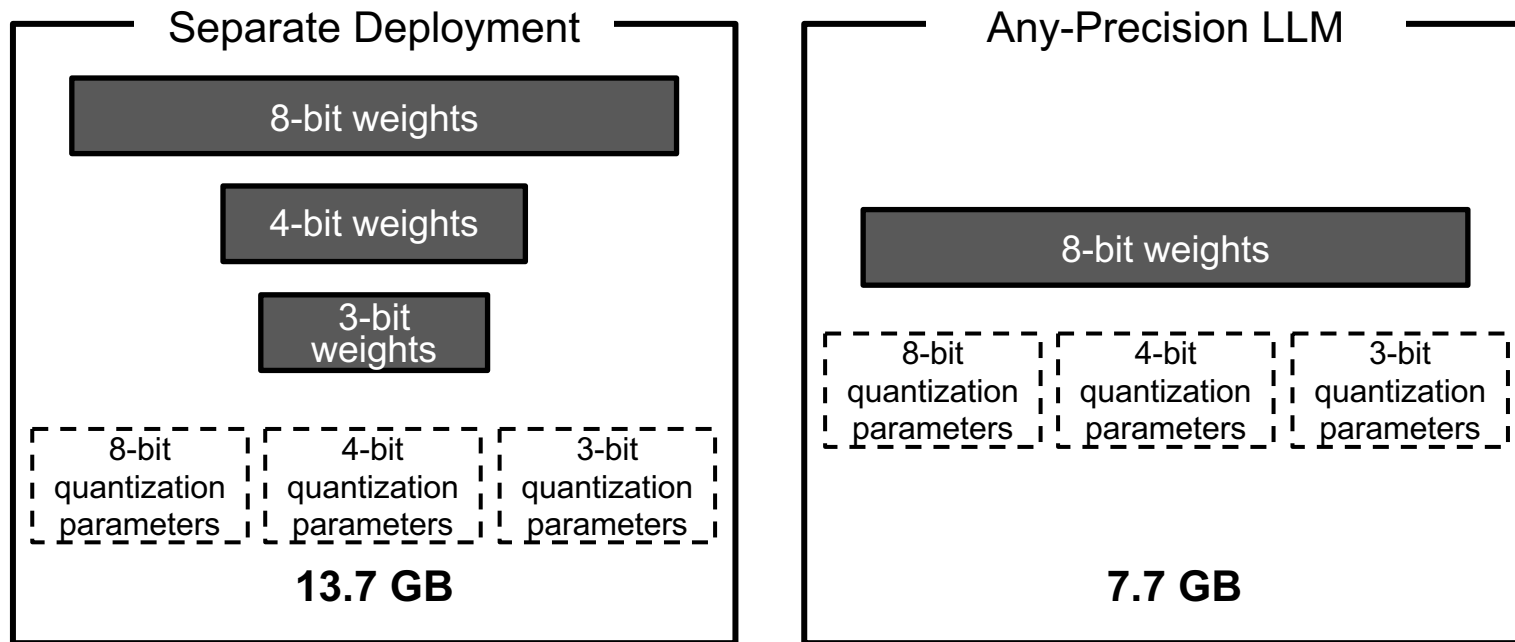


**1.49X saving**



# Potential Memory Saving of Any-Precision LLM (Llama-2-7B)

Supported bit-widths: {3, 4, 8}

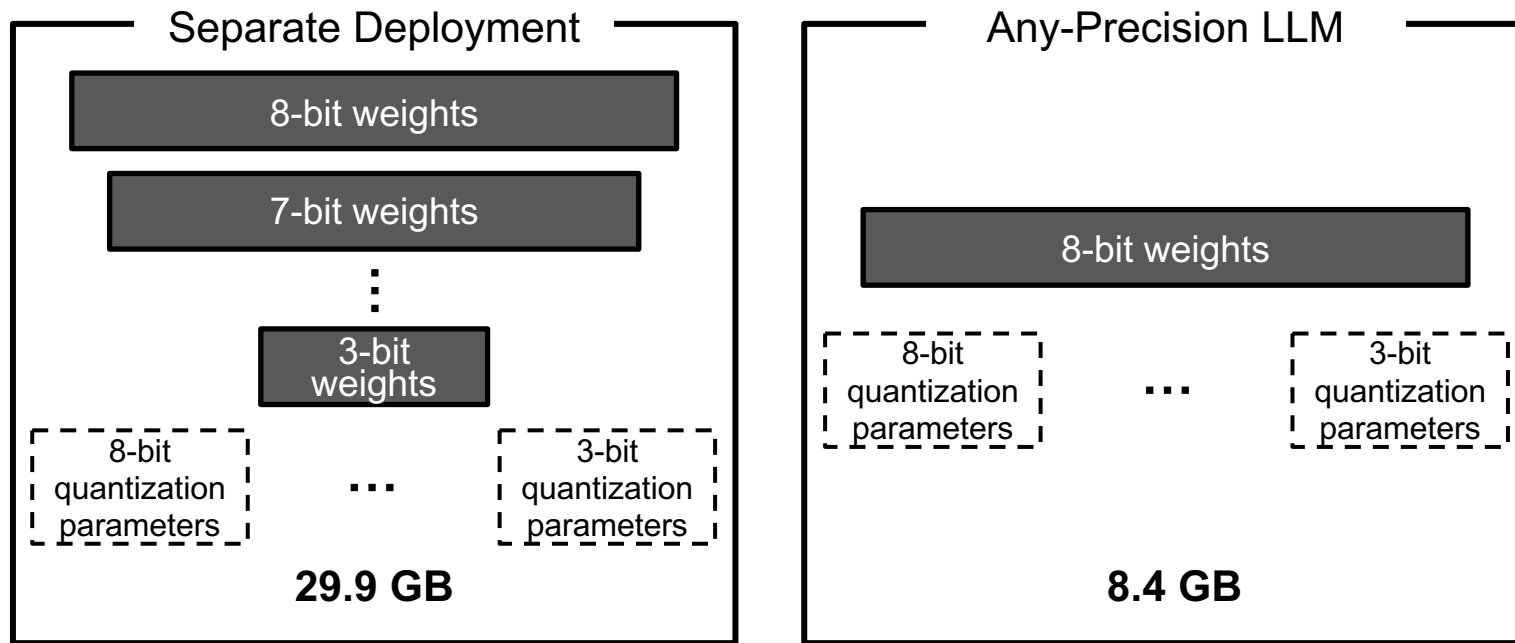


**1.76X saving**



# Potential Memory Saving of Any-Precision LLM (Llama-2-7B)

Supported bit-widths: {3, 4, 5, 6, 8}



**3.56X saving**



# Challenges of Any-Precision LLM

## ① High Training Cost

Original work adopts QAT (quantization-aware training) scheme for any-precision quantization


Compute Cluster



Energy Consumption



Dataset




Not affordable to most users

# Challenges of Any-Precision LLM

## ② Memory Bandwidth Saving

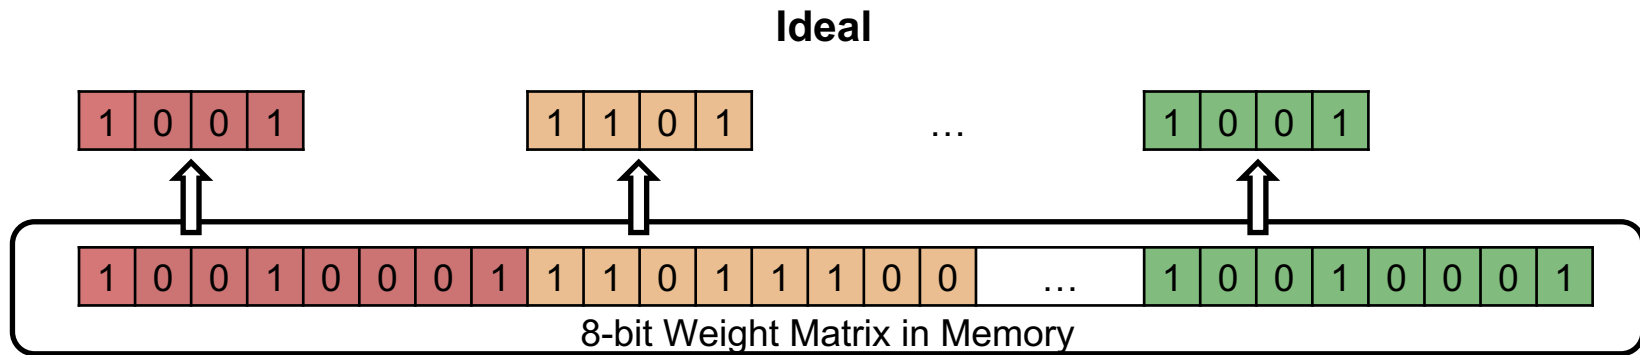
Original work does not save any memory bandwidth, resulting in no performance improvement for LLM



# Challenges of Any-Precision LLM

## ② Memory Bandwidth Saving

Original work does not save any memory bandwidth, resulting in no performance improvement for LLM






# Challenges of Any-Precision LLM

## ② Memory Bandwidth Saving

Original work does not save any memory bandwidth, resulting in no performance improvement for LLM



# Low-Cost Deployment of Multiple, Different-Sized LLMs

We make a strong case for any-precision quantization of LLM that **does not require training** and leads to **real end-to-end inference speed-ups**.

- **Lightweight method for any-precision quantization of LLMs** leveraging post-training quantization (PTQ) framework, called incremental upscaling
- **Specialized software engine** for efficient serving of any-precision LLM adopting a bitplane-based memory layout



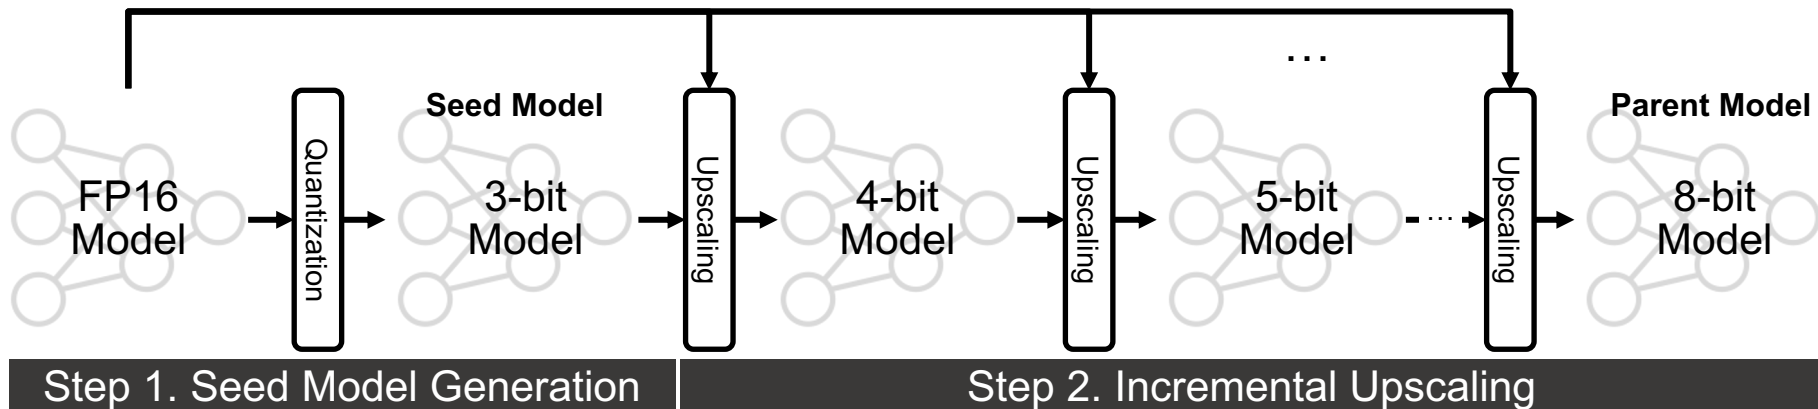
# Low-Cost Deployment of Multiple, Different-Sized LLMs

We make a strong case for any-precision quantization of LLM that **does not require training** and leads to **real end-to-end inference speed-ups**.

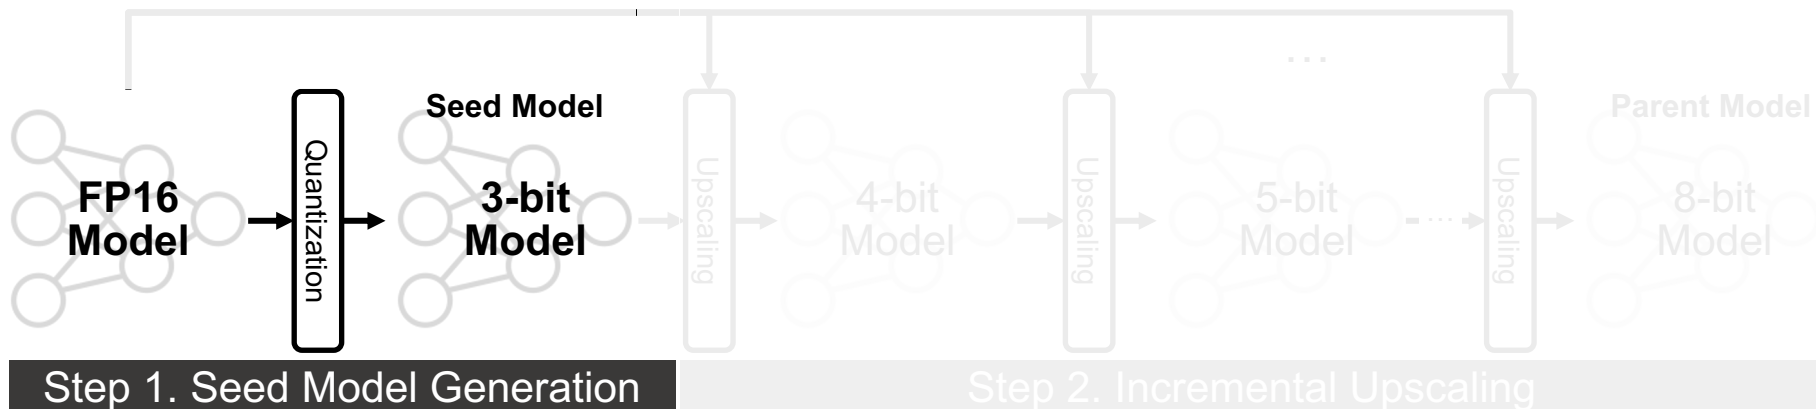
- **Lightweight method for any-precision quantization of LLMs** leveraging post-training quantization (PTQ) framework, called incremental upscaling
- **Specialized software engine** for efficient serving of any-precision LLM adopting a bitplane-based memory layout



# Any-Precision Quantization with Incremental Upscaling





# Any-Precision Quantization with Incremental Upscaling



$$\begin{pmatrix} 0.58 \\ 0.47 \\ \dots \\ -0.12 \end{pmatrix} \rightarrow \begin{pmatrix} 101 \\ 101 \\ \dots \\ 010 \end{pmatrix}$$


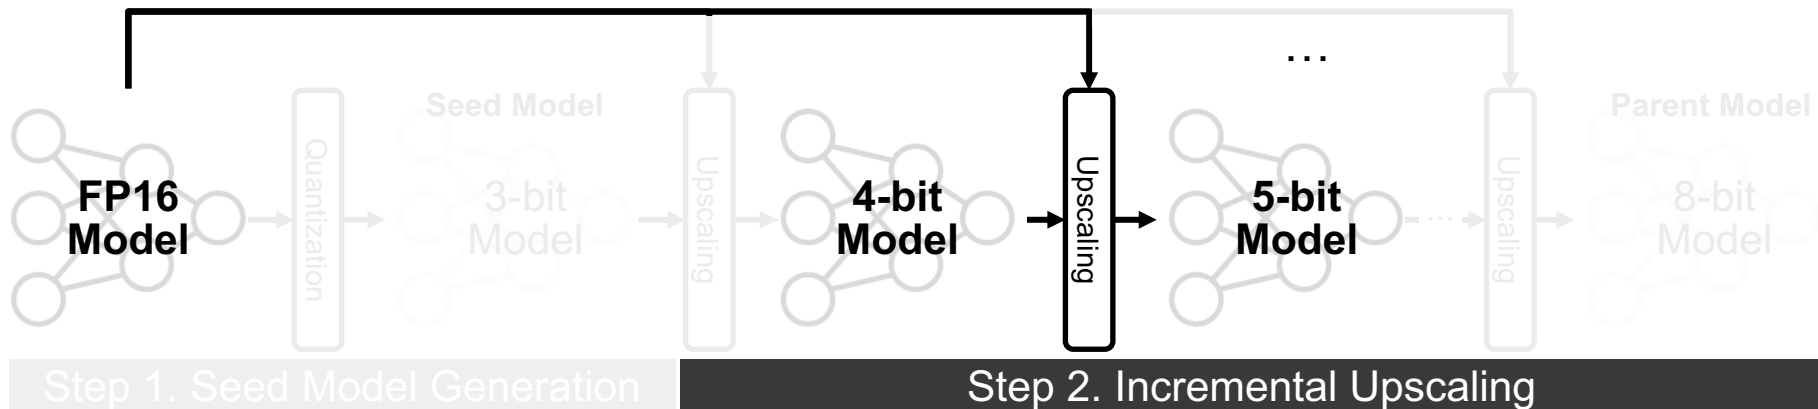
Quantize the model to the minimum supported bit-width (3-bit)

# Any-Precision Quantization with Incremental Upscaling




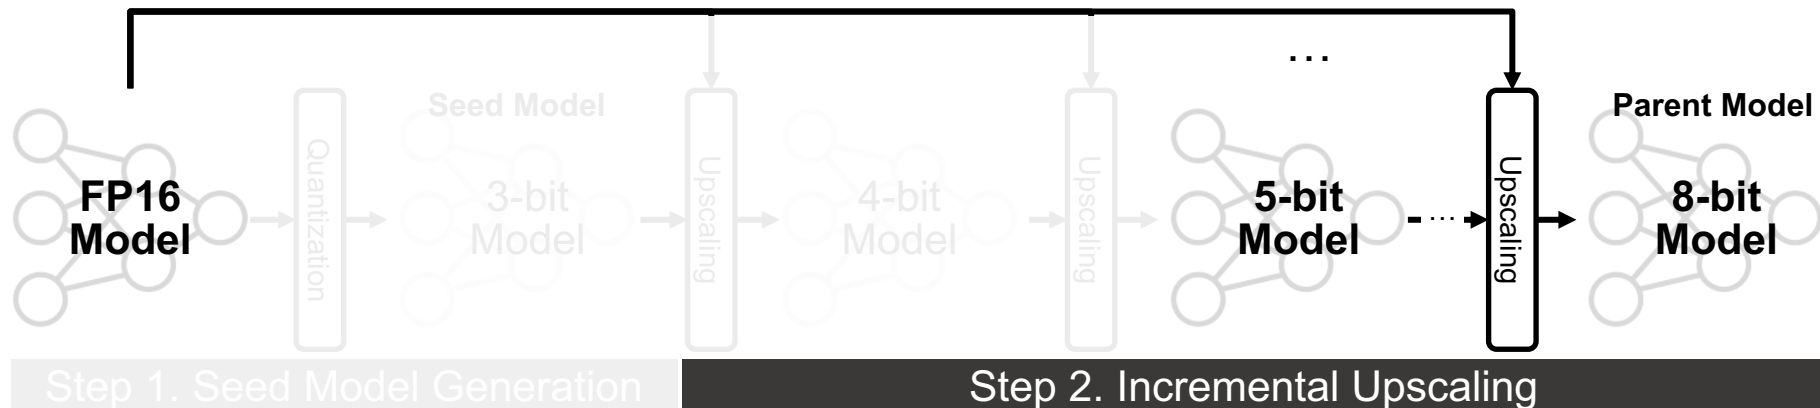
Upscale 3-bit model to 4-bit model

# Any-Precision Quantization with Incremental Upscaling



Upscale 4-bit model to 5-bit model

# Any-Precision Quantization with Incremental Upscaling



Upscale upto 8-bit model (parent model)



# Selection of Backbone Quantization Method

- A particular quantization method must be adopted as a backbone for **seed model generation** and **incremental upscaling (IU)**



# Which Quantization Method to Use?

- A particular quantization method must be adopted as a backbone for **seed model generation** and **incremental upscaling (IU)**



## **Requirement #1. Low-Cost**

No training; PTQ (post-training quantization) methods

# Which Quantization Method to Use?

- A particular quantization method must be adopted as a backbone for **seed model generation** and **incremental upscaling (IU)**



## **Requirement #1. Low-Cost**

No training; PTQ (post-training quantization) methods



## **Requirement #2. High-Performance**

Demonstrate SOTA results in terms of PPL/downstream task evaluation

# Which Quantization Method to Use?

- A particular quantization method must be adopted as a backbone for **seed model generation** and **incremental upscaling (IU)**



## **Requirement #1. Low-Cost**

No training; PTQ (post-training quantization) methods



## **Requirement #2. High-Performance**

Demonstrate SOTA results in terms of PPL/downstream task evaluation




## **Requirement #3. Compatible with IU**

Easily extended to support IU

# Incremental Upscaling with SqueezeLLM

\* SqueezeLLM: **non-uniform** quantization by **weighted K-means clustering**



Divide each cluster into two sub-clusters by **weighted K-means clustering**

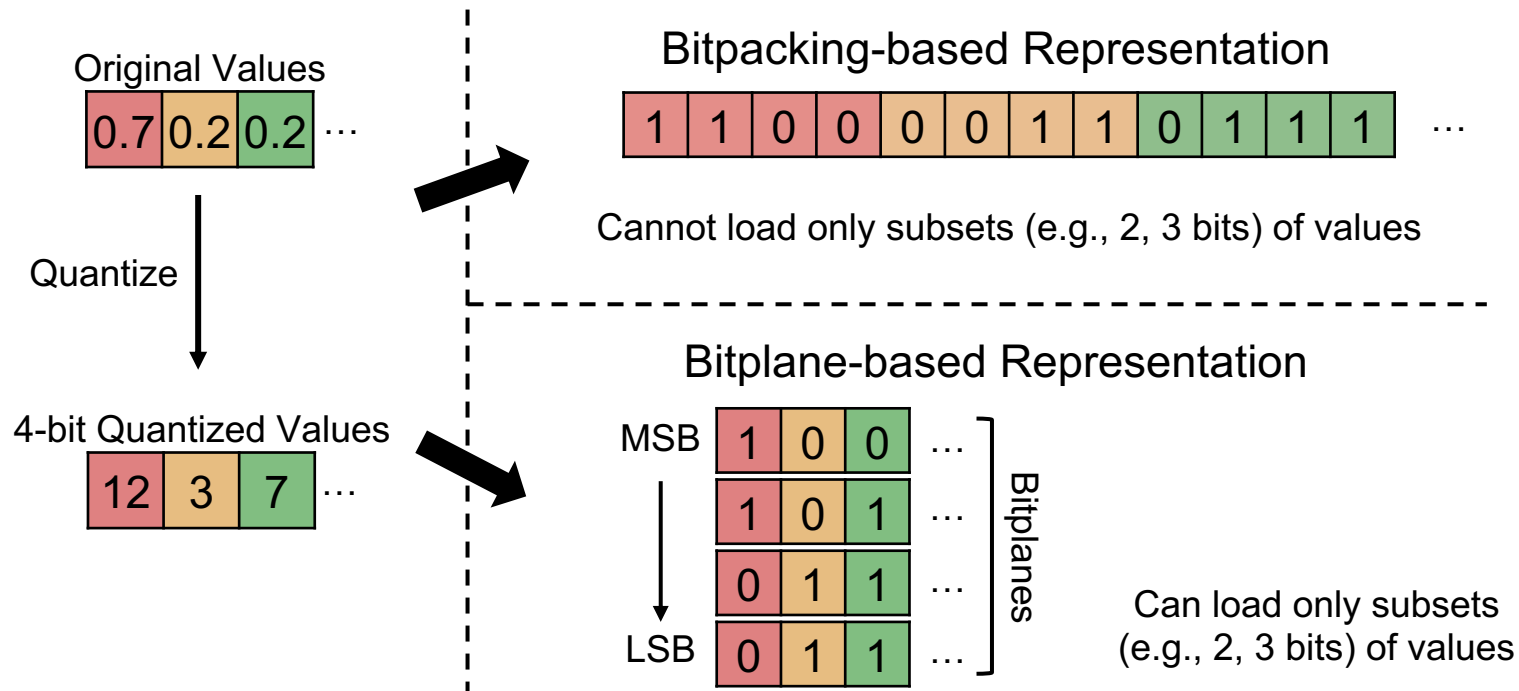
# Low-Cost Deployment of Multiple, Different-Sized LLMs

We make a strong case for any-precision quantization of LLM that **does not require training** and leads to a **real end-to-end inference speed-up**.

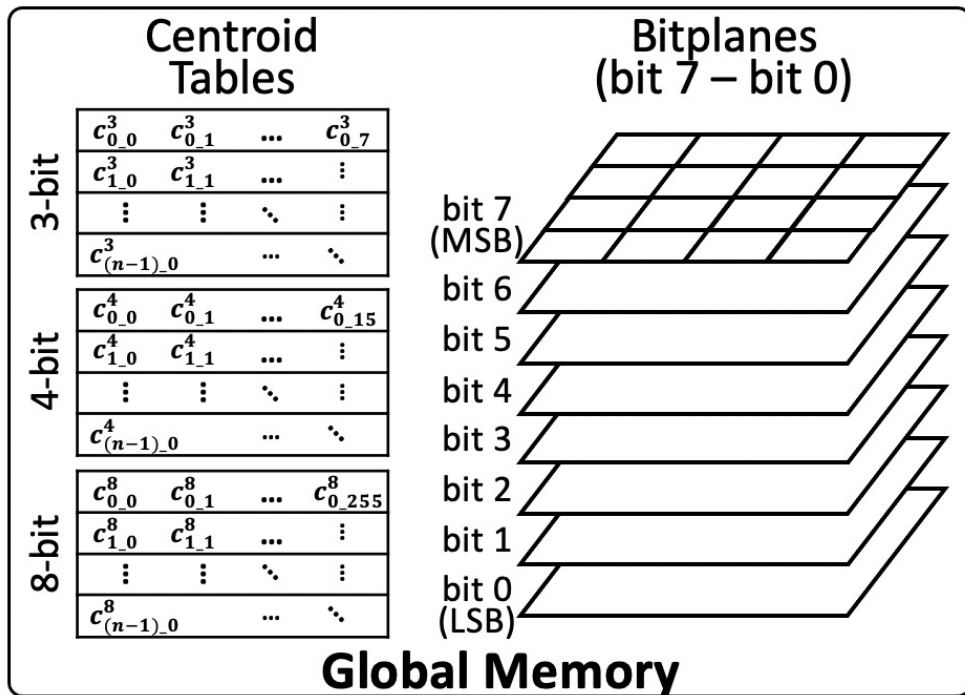
- **Lightweight method for any-precision quantization of LLMs** leveraging post-training quantization (PTQ) framework, called incremental upscaling
- **Specialized software engine** for efficient serving of any-precision LLM adopting a bitplane-based memory layout



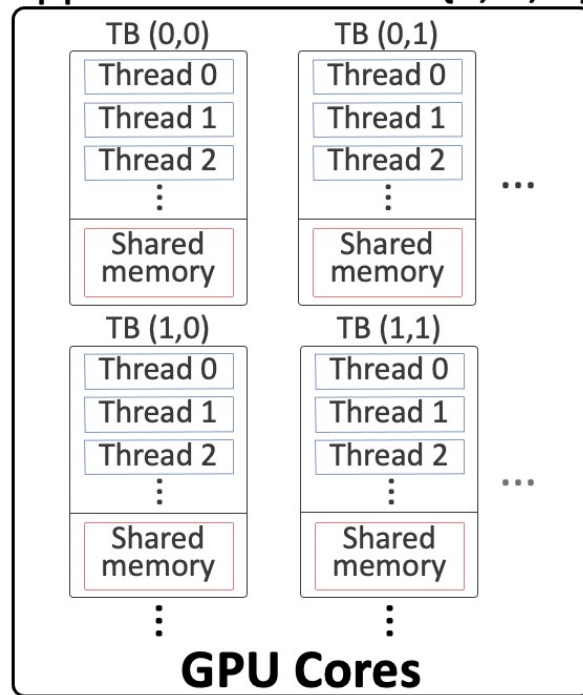
# Bitpacking-based vs Bitplane-based



# Overview of Our Specialized Engine

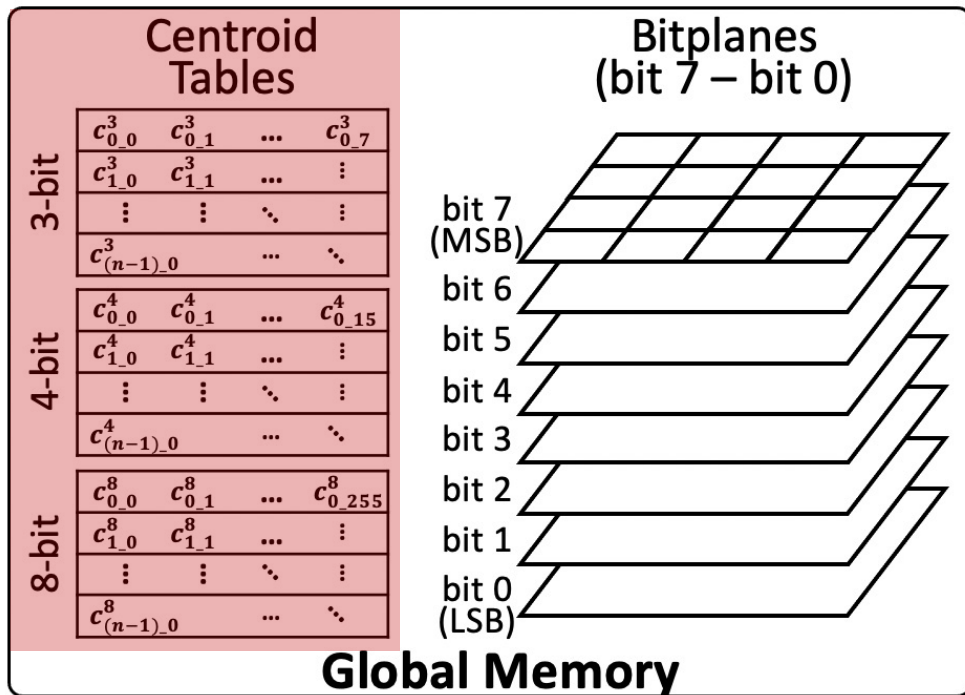


Supported bit-widths: {3, 4, 8}

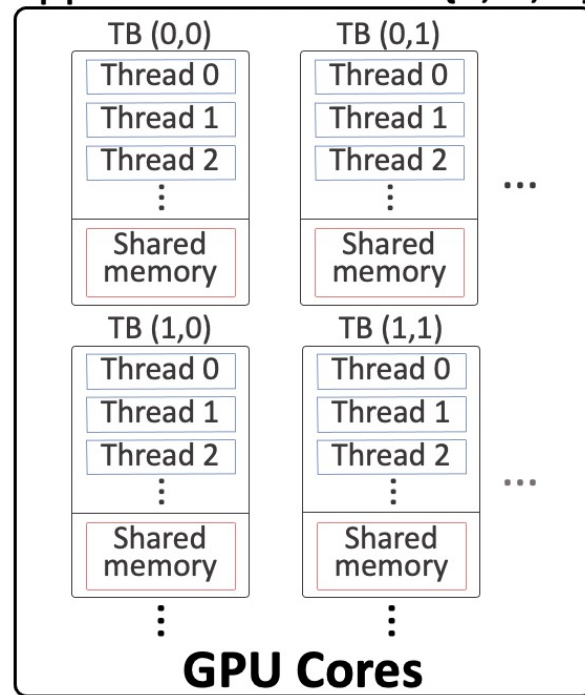




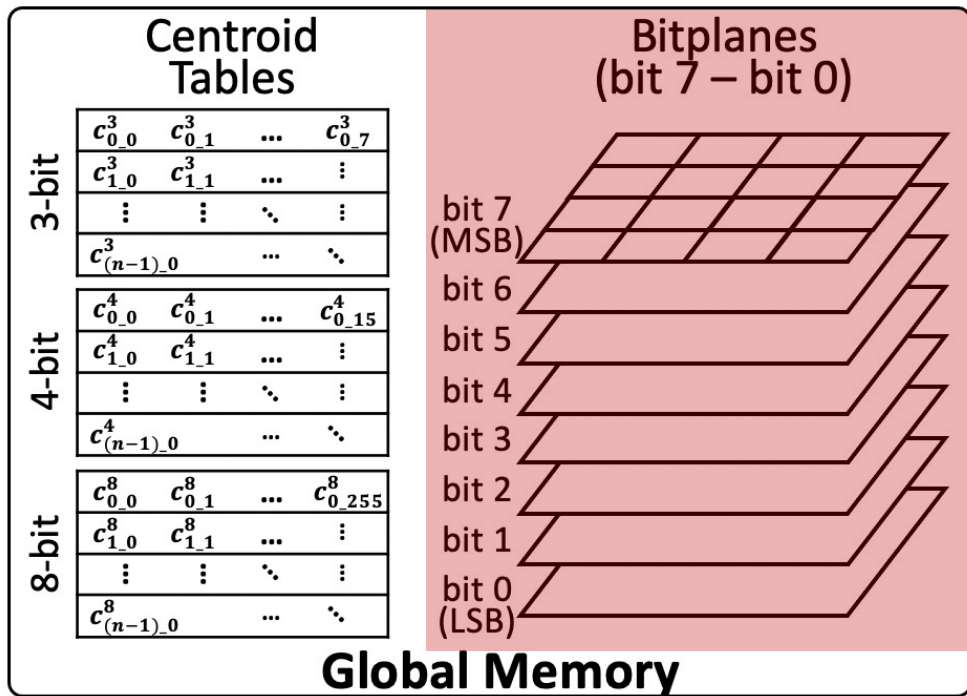
# Overview of Our Specialized Engine



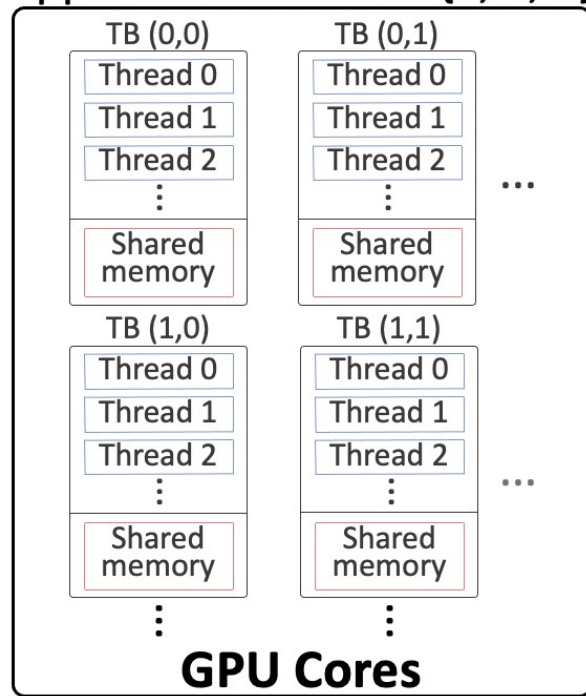
Supported bit-widths: {3, 4, 8}



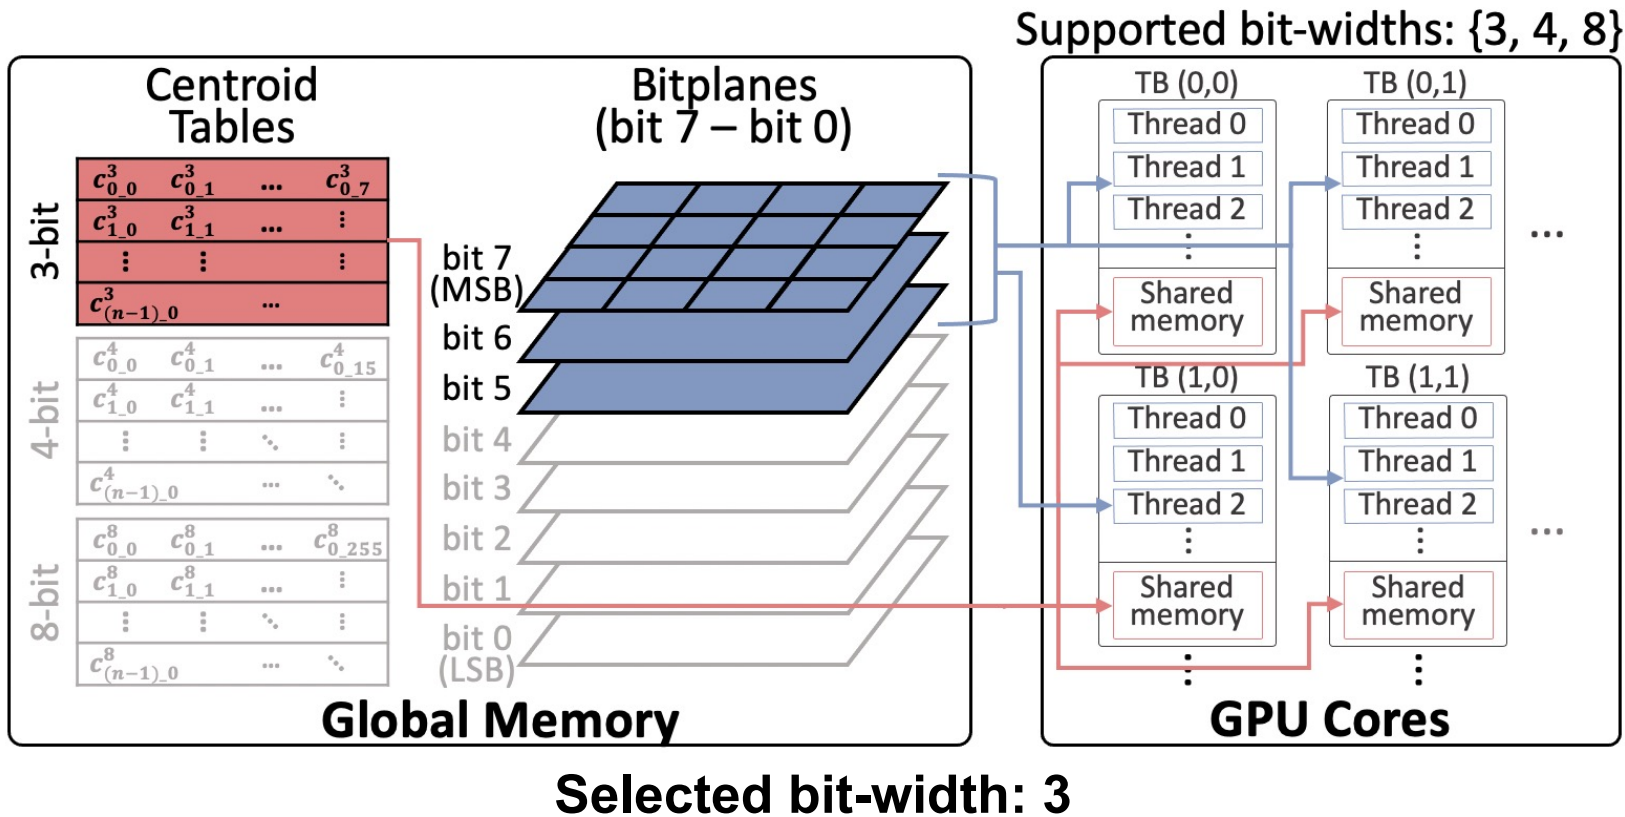
# Overview of Our Specialized Engine



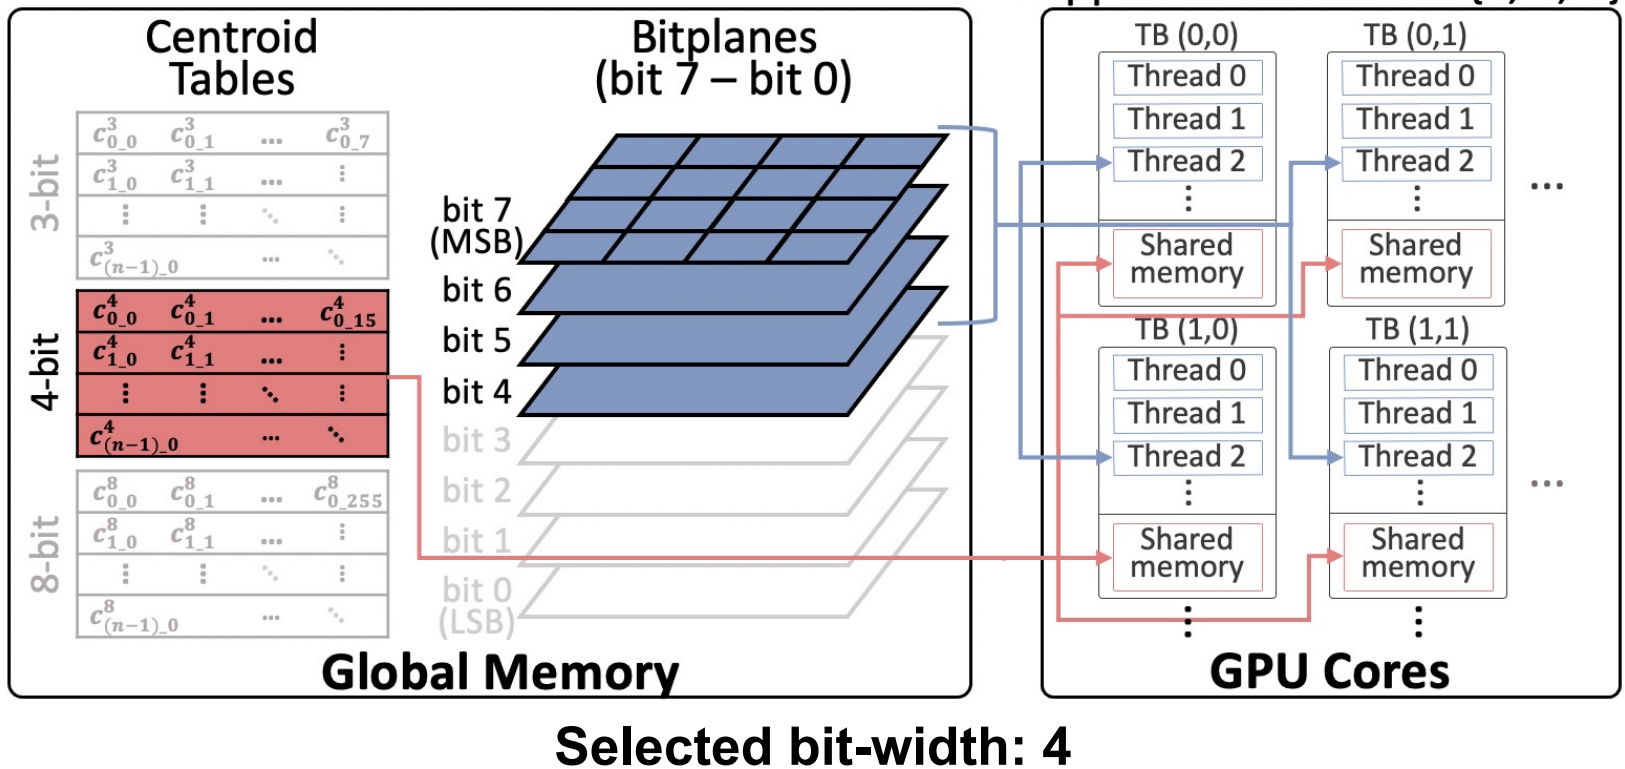
Supported bit-widths: {3, 4, 8}



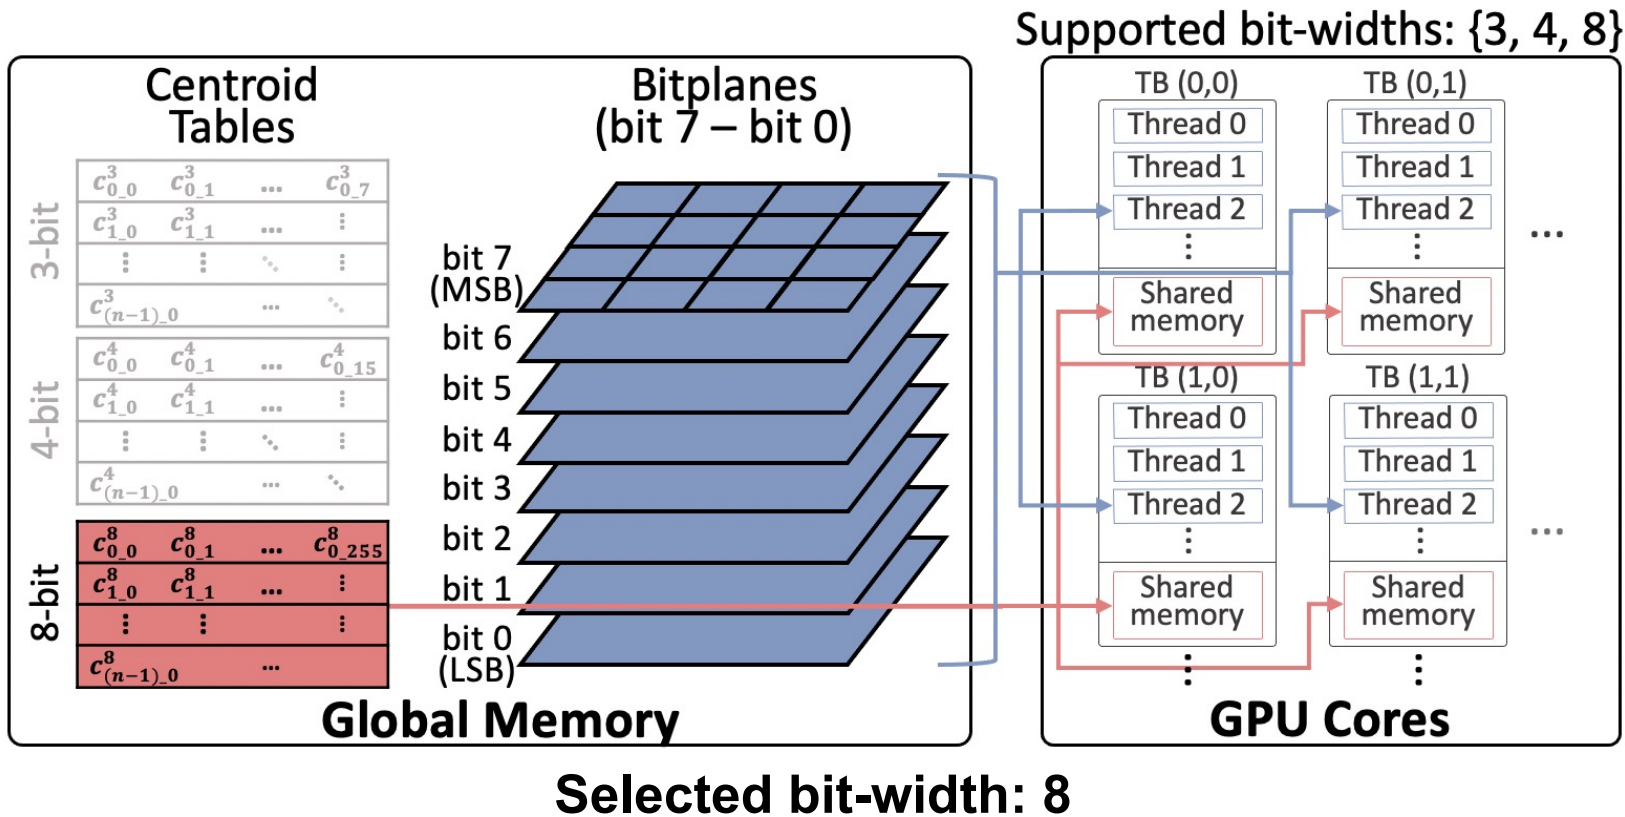
# Overview of Our Specialized Engine




# Overview of Our Specialized Engine



# Overview of Our Specialized Engine




# Thread-Level Operations




Five step thread-level operations assuming a bit-width of 4

# Thread-Level Operations



**Load input activations and weights from global memory**


# Thread-Level Operations



Align the bits of each weight contiguously  
(bit-level transpose)




# Thread-Level Operations



Get addresses for centroid table lookup by


- 1) Shift & Mask
- 2) Add offset (base address)

# Thread-Level Operations




## Fetch centroids and perform MAC

# GPU Kernel Optimization




# GPU Kernel Optimization



## Challenge ①

Uncoalesced memory access

# GPU Kernel Optimization




## Challenge ①

### Uncoalesced memory access

-> *Weight bitplane layout optimization*



# GPU Kernel Optimization



## Challenge ①


Uncoalesced memory access

-> *Weight bitplane layout optimization*

## Challenge ②

Bitwise operation overhead

# GPU Kernel Optimization



## Challenge ①

Uncoalesced memory access

-> *Weight bitplane layout optimization*

## Challenge ②

Bitwise operation overhead

-> *Efficient bit-transpose*

-> *Table lookup merge*



# Evaluation





# Evaluation

- **Key Result #1:**

A set of quantized models generated from incremental upscaling match the SOTA quantization results

- **Key Result #2:**

Our engine matches or even outperforms existing engines while providing memory-efficient any-precision support



# Evaluation

- **Key Result #1:**

A set of quantized models generated from incremental upscaling match the SOTA quantization results

- **Key Result #2:**

Our engine matches or even outperforms existing engines while providing memory-efficient any-precision support



# Perplexity on Wikitext-2

		3-bit	4-bit	5-bit	6-bit	7-bit	8-bit	FP16
Llama-2-7B	SqLLM	6.13	5.61	5.50	5.47	5.47	5.47	5.47
	SqLLM+IU		5.62	5.50	5.47	5.47	5.47	
	$\Delta$		$\uparrow$ 0.01	-	-	-	-	

Upscaled models (**SqLLM+IU**) match independently quantized models (**SqLLM**)

Similar trends observed for *other datasets* (PTB, C4) and *other models* (Mistral, OPT)



# Zero-shot Task Accuracy

\* Average accuracy across five tasks (ARC-easy/challenge, HellaSwag, PIQA, WinoGrande)

		3-bit	4-bit	5-bit	6-bit	7-bit	8-bit	FP16
Llama-2-7B	SqLLM	66.2	68.3	68.6	68.8	68.9	68.9	69.0
	SqLLM+IU		68.2	68.8	68.9	68.9	69.0	
	$\Delta$		↓ 0.01	↑ 0.02	↑ 0.01	-	↑ 0.01	

Upscaled models (**SqLLM+IU**) match independently quantized models (**SqLLM**)

Similar trends observed for *other models (Mistral, OPT)*



# Evaluation

- **Key Result #1:**

A set of quantized models generated from incremental upscaling match the SOTA quantization results




- **Key Result #2:**

Our engine matches or even outperforms existing engines while providing memory-efficient any-precision support



# Kernel Microbenchmark (GEMV)


\* Speedup over cuBLAS FP16 baseline

			(1, 4096) x (4096, 4096)						(1, 11008) x (4096, 11008)					
			3-bit	4-bit	5-bit	6-bit	7-bit	8-bit	3-bit	4-bit	5-bit	6-bit	7-bit	8-bit
	RTX 4090	Ours	3.99×	3.03×	2.61×	2.18×	1.87×	1.56×	4.45×	3.42×	2.74×	2.28×	2.08×	1.81×
		SqLLM	3.69×	3.07×	-	-	-	-	4.22×	3.17×	-	-	-	-
	RTX 4070 Laptop	Ours	4.97×	3.73×	3.01×	2.51×	2.10×	1.76×	5.29×	3.66×	3.05×	2.52×	2.13×	1.87×
		SqLLM	4.74×	3.66×	-	-	-	-	5.29×	3.93×	-	-	-	-
	Jetson AGX Orin	Ours	3.84×	3.02×	2.56×	2.33×	2.10×	1.78×	4.35×	2.96×	2.54×	2.52×	2.16×	1.86×
		SqLLM	3.15×	1.94×	-	-	-	-	3.36×	2.04×	-	-	-	-

Our kernel outperforms for the most cases even with any-precision support



# End-to-End Speedup



Generation throughput of Llama-2-7B using our kernel integrated with TensorRT-LLM



# What's Next?





# Next Step: Any-Precision LLM for Datacenter Settings

Single Batch



Large Batch

Memory-Bound



Compute-Bound


Custom Number  
Formats



FP16, FP8, FP4



# Next Step: Any-Precision LLM for Datacenter Settings



# Summary

Any-precision LLM is a memory-efficient and cost-effective solution for deployment of multiple, different sized LLMs.

- **Reduce memory cost** of deployment of multiple LLMs by any-precision quantization
- Propose **lightweight any-precision quantization scheme** for LLMs which produces quantized models with SOTA results
- Propose **specialized software engine** for efficient serving of any-precision LLMs
- Currently exploring ways to extend this concept to datacenter settings



Paper



Code

