

# RiF: Improving Read Performance of Modern SSDs Using an On-Die Early-Retry Engine

Myoungjun Chun<sup>1</sup>, Jaeyong Lee<sup>1</sup>, Myungsuk Kim<sup>2</sup>,  
Jisung Park<sup>3</sup>, and Jihong Kim<sup>1</sup>

---

<sup>1</sup>Seoul National University

<sup>2</sup>Kyungpook National University

<sup>3</sup>POSTECH

# Talk Outline

001

Read Retry in Modern SSDs

002

Limitation of the Existing Read Retry Optimization Techniques






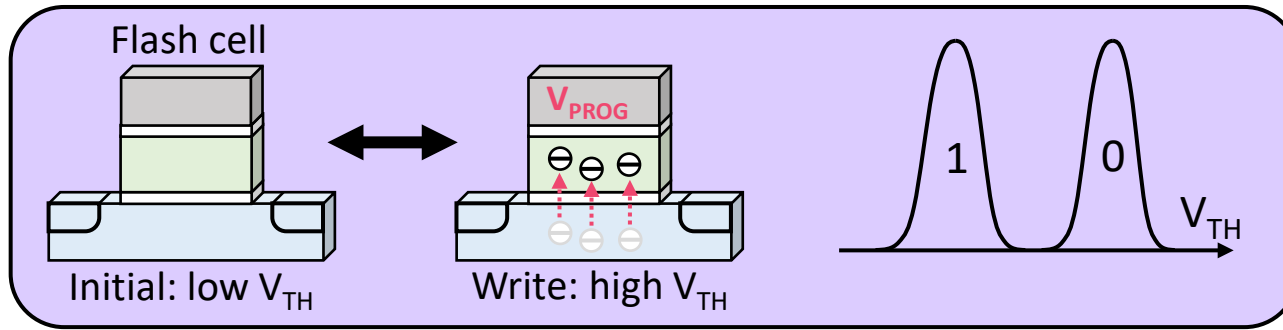
003

RiF: Retry-in-Flash

004

Evaluation Results and Conclusion

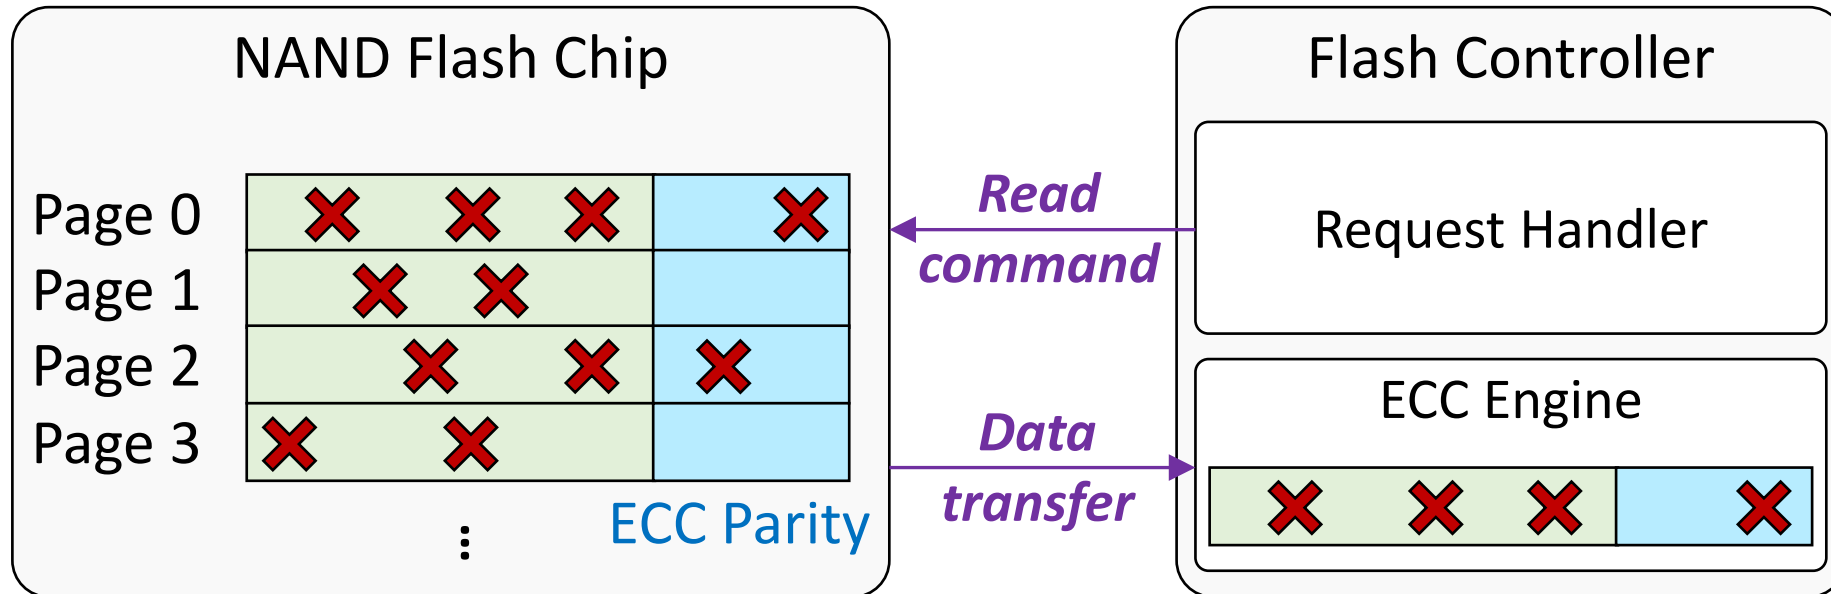
# Errors in NAND Flash Memory



High-density NAND flash memory is *highly error-prone*

# Error-Correcting Codes (ECC)

- Store **redundant information (ECC parity)** for error correction




When # of raw bit errors (RBER) exceeds ECC correction capability,  
**uncorrectable errors** occur

# Read-Retry in NAND Flash Memory

- Read-retry **re-reads** the target page with adjusted  $V_{\text{REF}}$  values


ECC  
engine

*ECC capability = 80 bits*



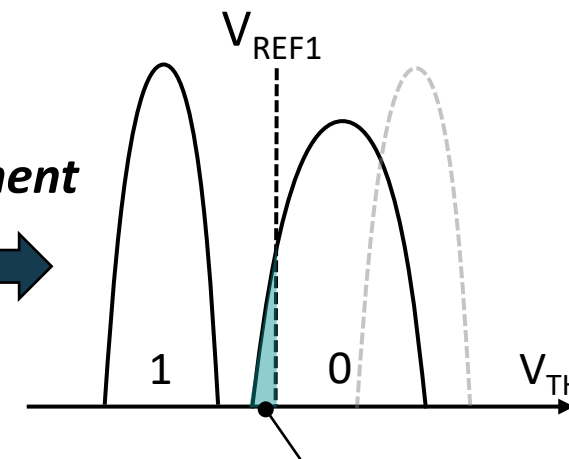
*# of raw bit errors = 150*

$V_{\text{REF}}$  adjustment



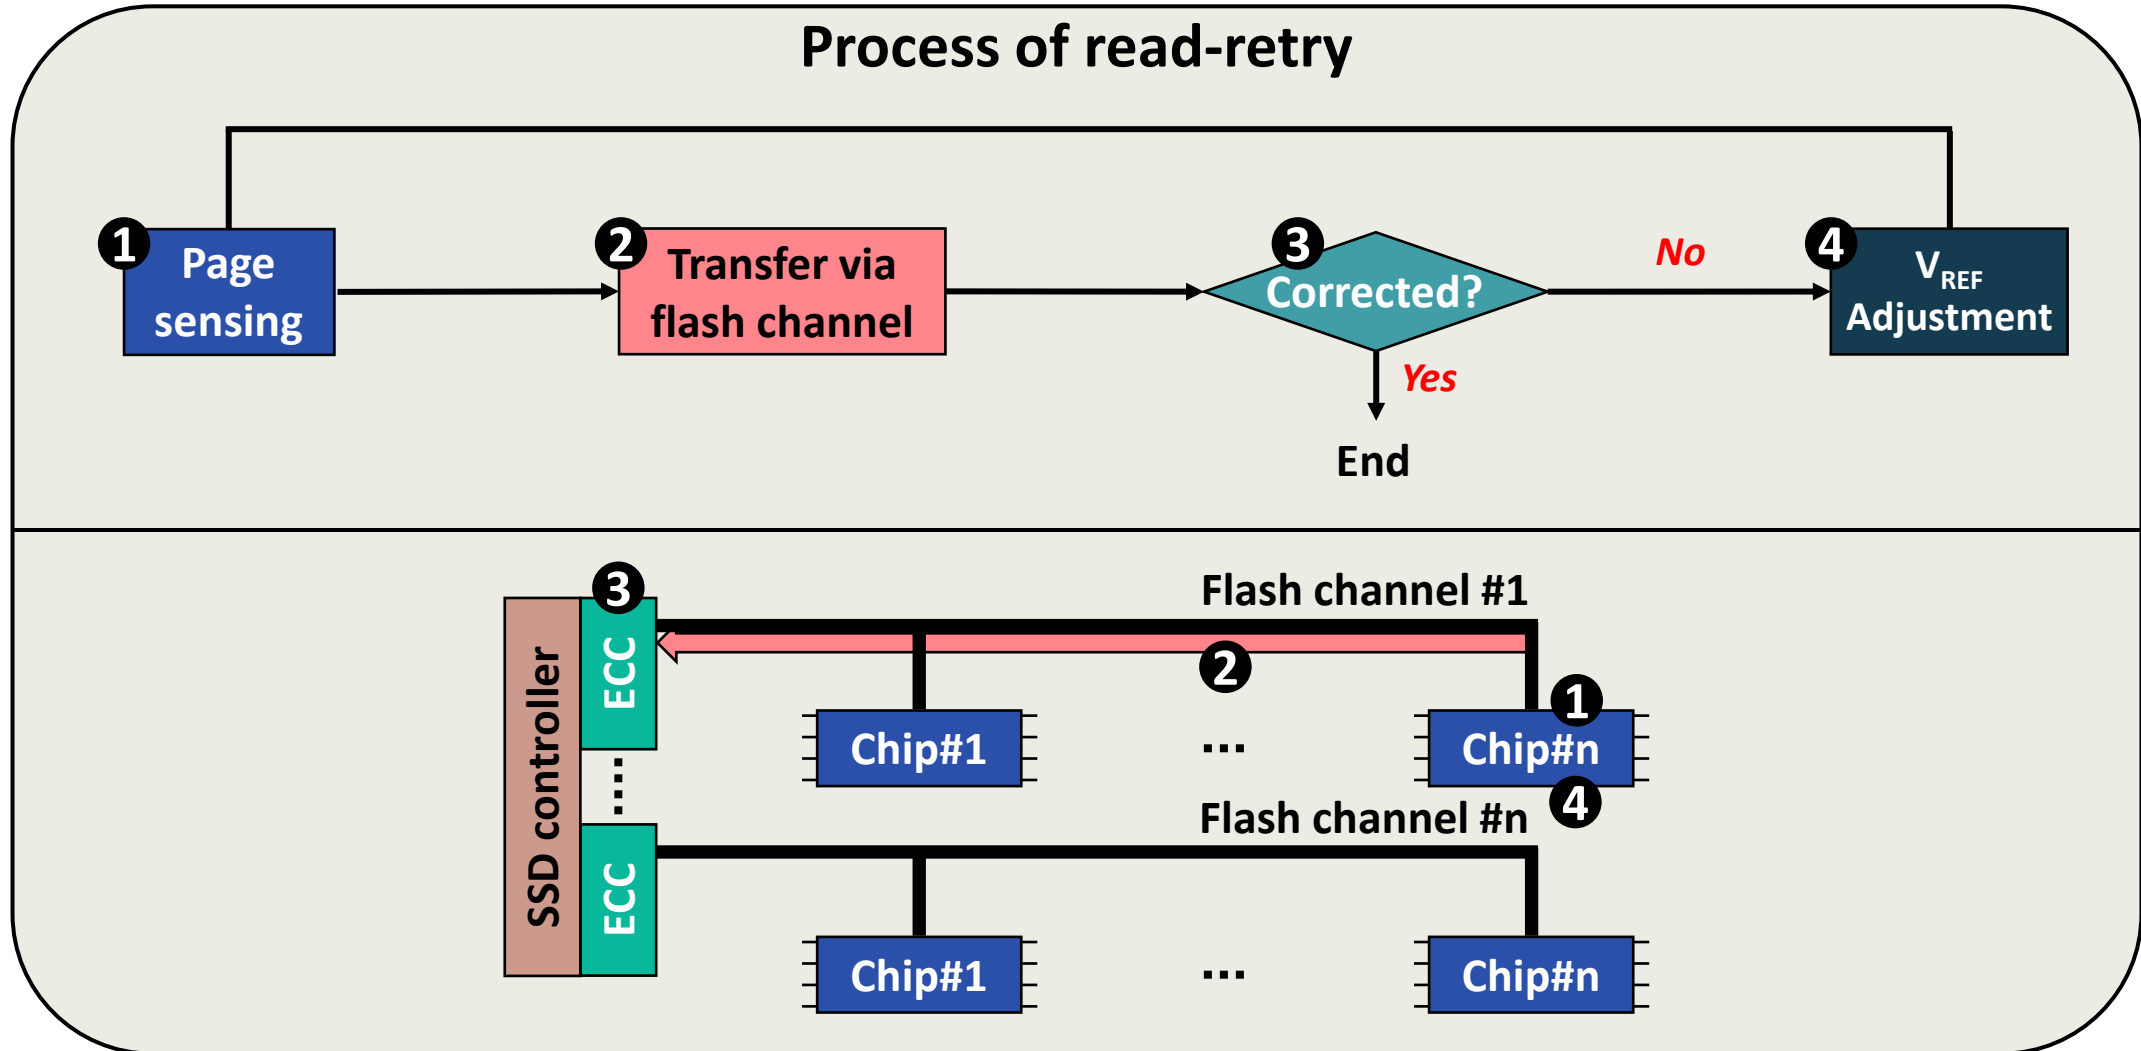
*# of raw bit errors = 100*

$V_{\text{REF}}$  adjustment




*# of raw bit errors = 60*

# Read-Retry in NAND Flash Memory (cont'd)



# Frequency of Read-Retry in Modern NAND Flash

Characterization results using 160 real chips



Read-retry (i.e., ECC failure) *occurs very frequently* in modern NAND flash

# Talk Outline

001

Read Retry in Modern SSDs

002

Limitation of the Existing Read Retry Optimization Techniques

003

RiF: Retry-in-Flash

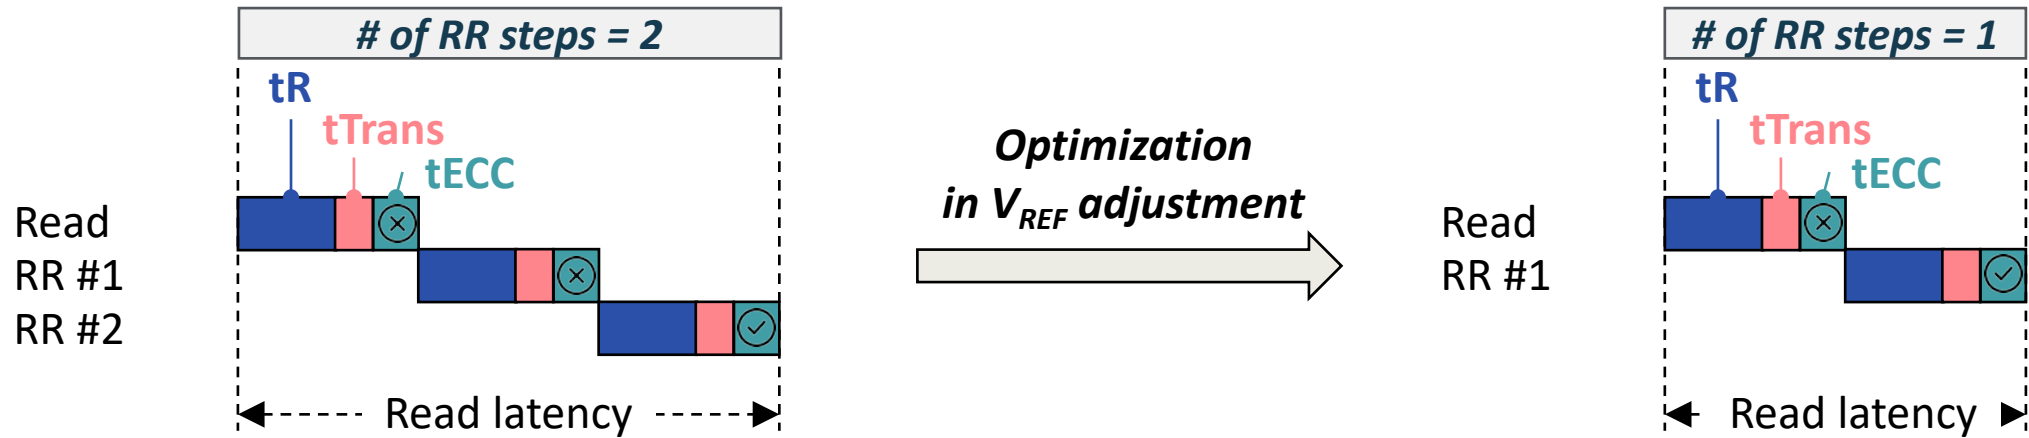
004

Evaluation Results and Conclusion




# Read-Retry Optimization

- To mitigate the performance overhead of frequent read-retry, many prior works focused on *reducing # of read-retry steps*




- State-of-the-art: Swift-Read (ISSCC22) *reduces # of read-retry steps to almost 1*



# Impact of Read-Retry on SSD's Performance

- Compare I/O bandwidth of two SSDs using MQSim-E
  - $SSD_{real}$ : An SSD w/ *optimal* read-retry solution (*# of RR steps = 1*)
  - $SSD_{ideal}$ : An ideal SSD w/ no read-retry

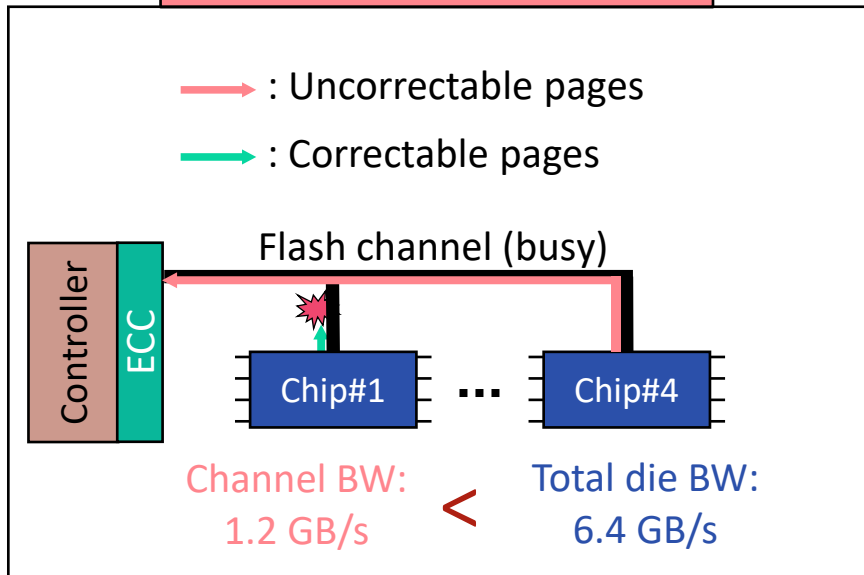


Read-Retry degrades SSD's performance significantly even  
*with the optimal solution*


# Root Cause Analysis

- **Root Cause:** wasted effort for *uncorrectable pages*

## Waste 1: page transfer



## Waste 2: long ECC decoding time



# Talk Outline

001

Read Retry in Modern SSDs

002

Limitation of the Existing Read Retry Optimization Techniques

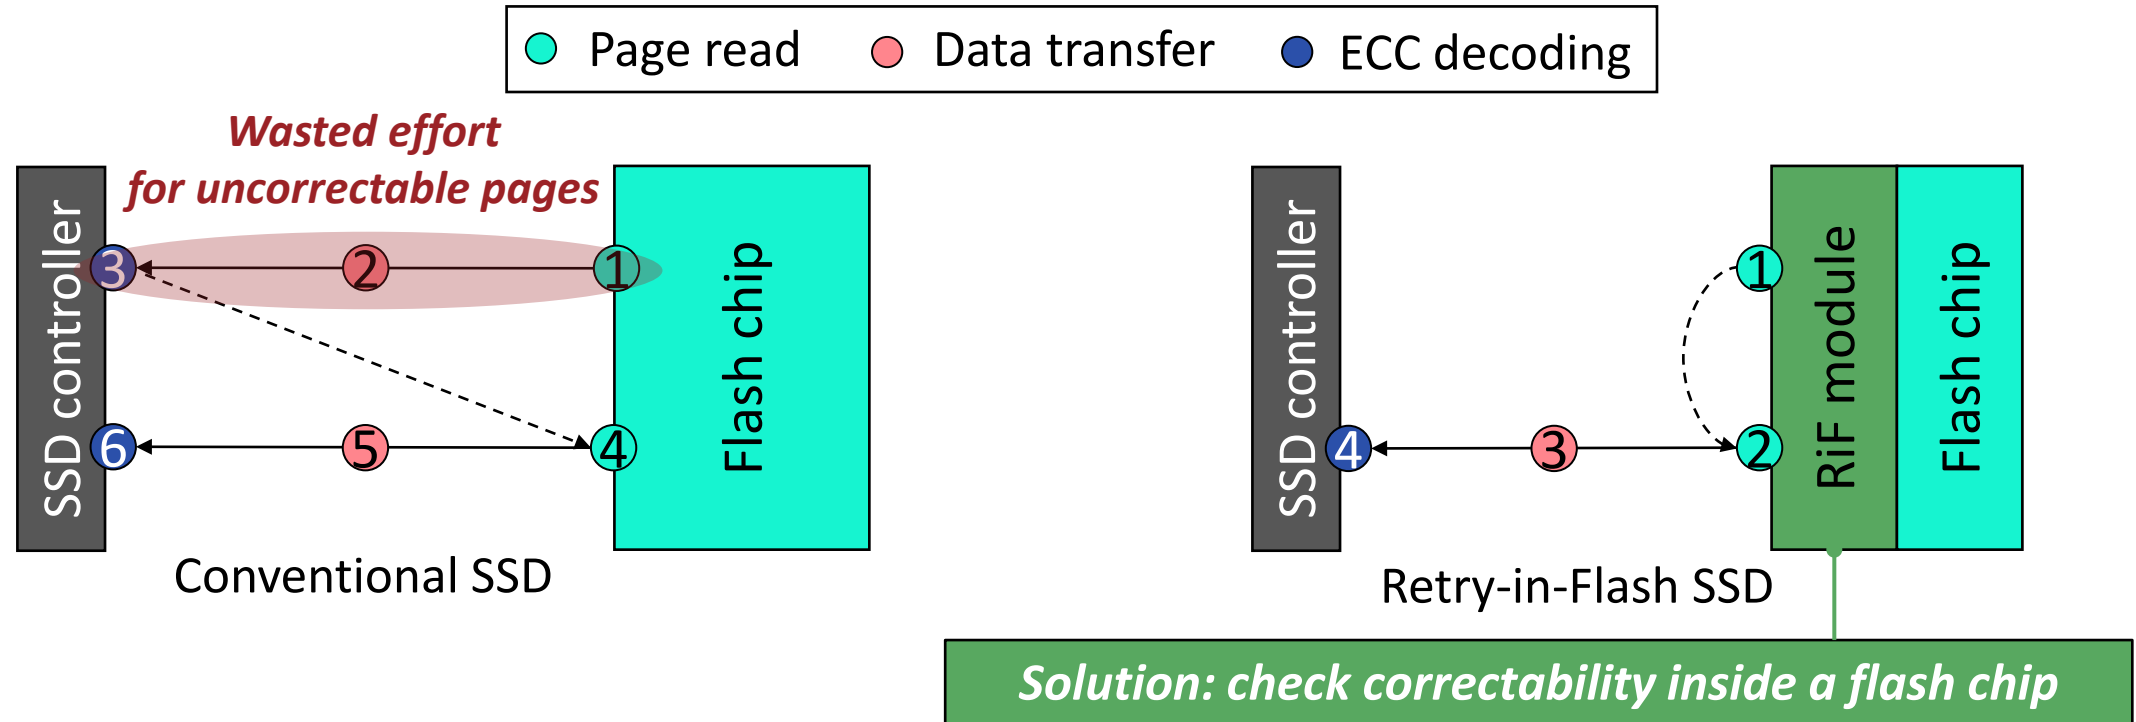
003

RiF: Retry-in-Flash

004

Evaluation Results and Conclusion

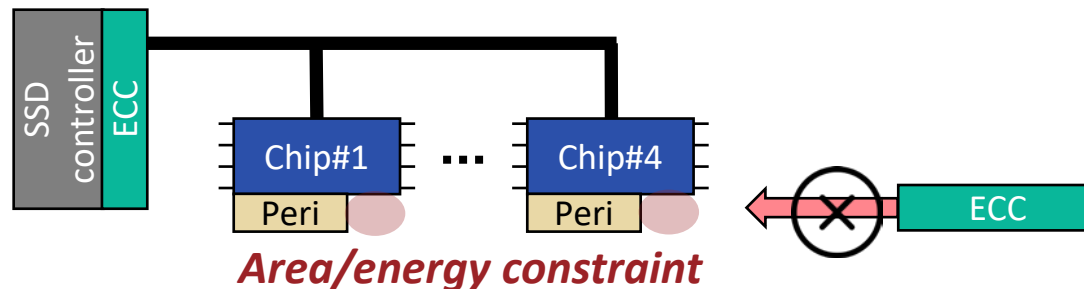
# Key Idea: Retry-in-Flash (RiF)



**Effect 1: No wasted page transfer (②)**  
**Effect 2: No wasted decoding time (③)**

# Implementation Challenge

- Hard constraints on *in-flash RiF module*



- **Optimization 1:** Subpage-based prediction



All data in pages must be **randomized** in modern NAND flash memory


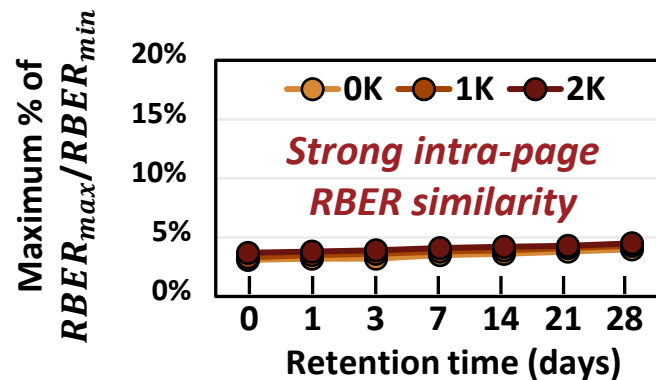
original data: 00000001

Randomizer

01010101


Chip

RB ER dist. of 4-KiB chunks in a 16-KiB page



# Implementation Challenge (cont'd)

- **Optimization 2:** Prediction-centric module



**Uncorrectable pages** do not require *entire ECC decoding*  
Instead, we simply want to know *if a page is correctable or not*

## Syndrome weight-based prediction

$H$  = a parity check matrix,  $\mathbf{c}$  = a codeword =  $(c_0, c_1 \dots c_{N-1})$ ,  $\mathbf{S}$  = a syndrome vector =  $(s_0, s_1, s_2, \dots s_{M-1})$

$$H = \begin{bmatrix} 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 0 \end{bmatrix}$$

$M=4$  (rows),  $N=8$  (columns)

1's represent the bit position of a codeword

$$\mathbf{S}^T = H\mathbf{c}^T = \begin{bmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{bmatrix} = \begin{bmatrix} c_1 \oplus c_3 \oplus c_4 \oplus c_7 \\ c_0 \oplus c_1 \oplus c_2 \oplus c_5 \\ c_2 \oplus c_5 \oplus c_6 \oplus c_7 \\ c_0 \oplus c_3 \oplus c_4 \oplus c_6 \end{bmatrix}$$

$s_3$  is 0 with no errors in  $c_0, c_3, c_4, c_6$

$s_3$  is 1 with any bit error

Syndrome weight  $SW$   
(i.e.,  $s_0 + s_1 + s_2 + s_3$ )

$SW > thr?$

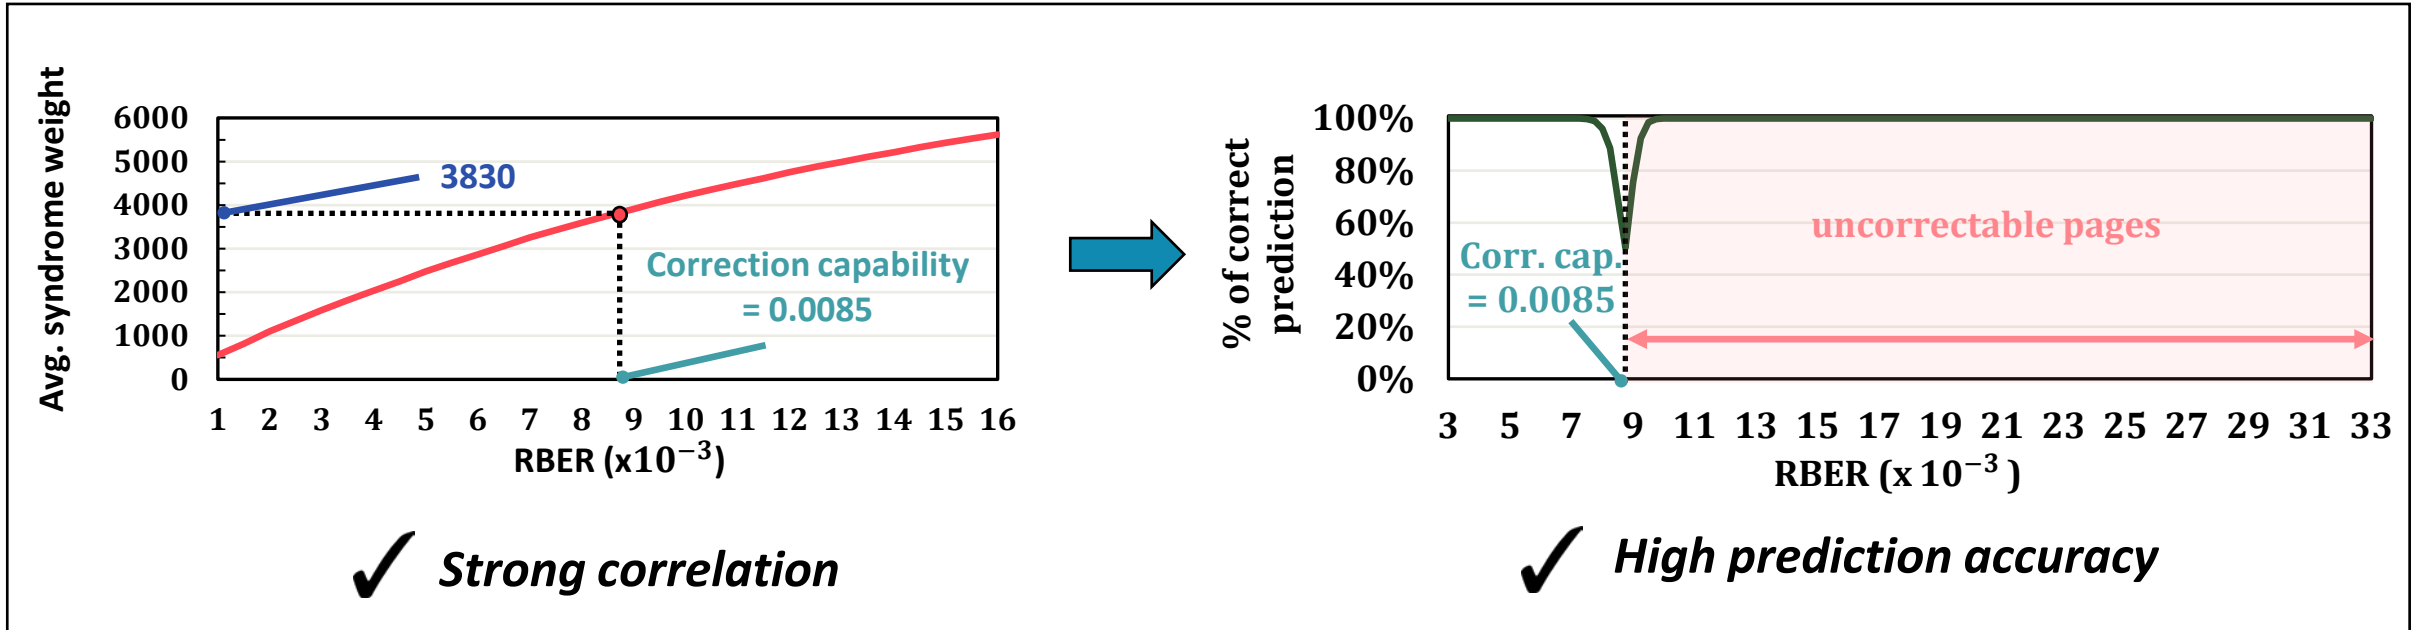
*Dout*

no

yes

*In-flash read-retry*

# Validation Results



RiF module achieves **98.7% prediction accuracy** based on the **syndrome weight-based prediction**



# Talk Outline

001

Read Retry in Modern SSDs

002

Limitation of the Existing Read Retry Optimization Techniques

003

RiF: Retry-in-Flash

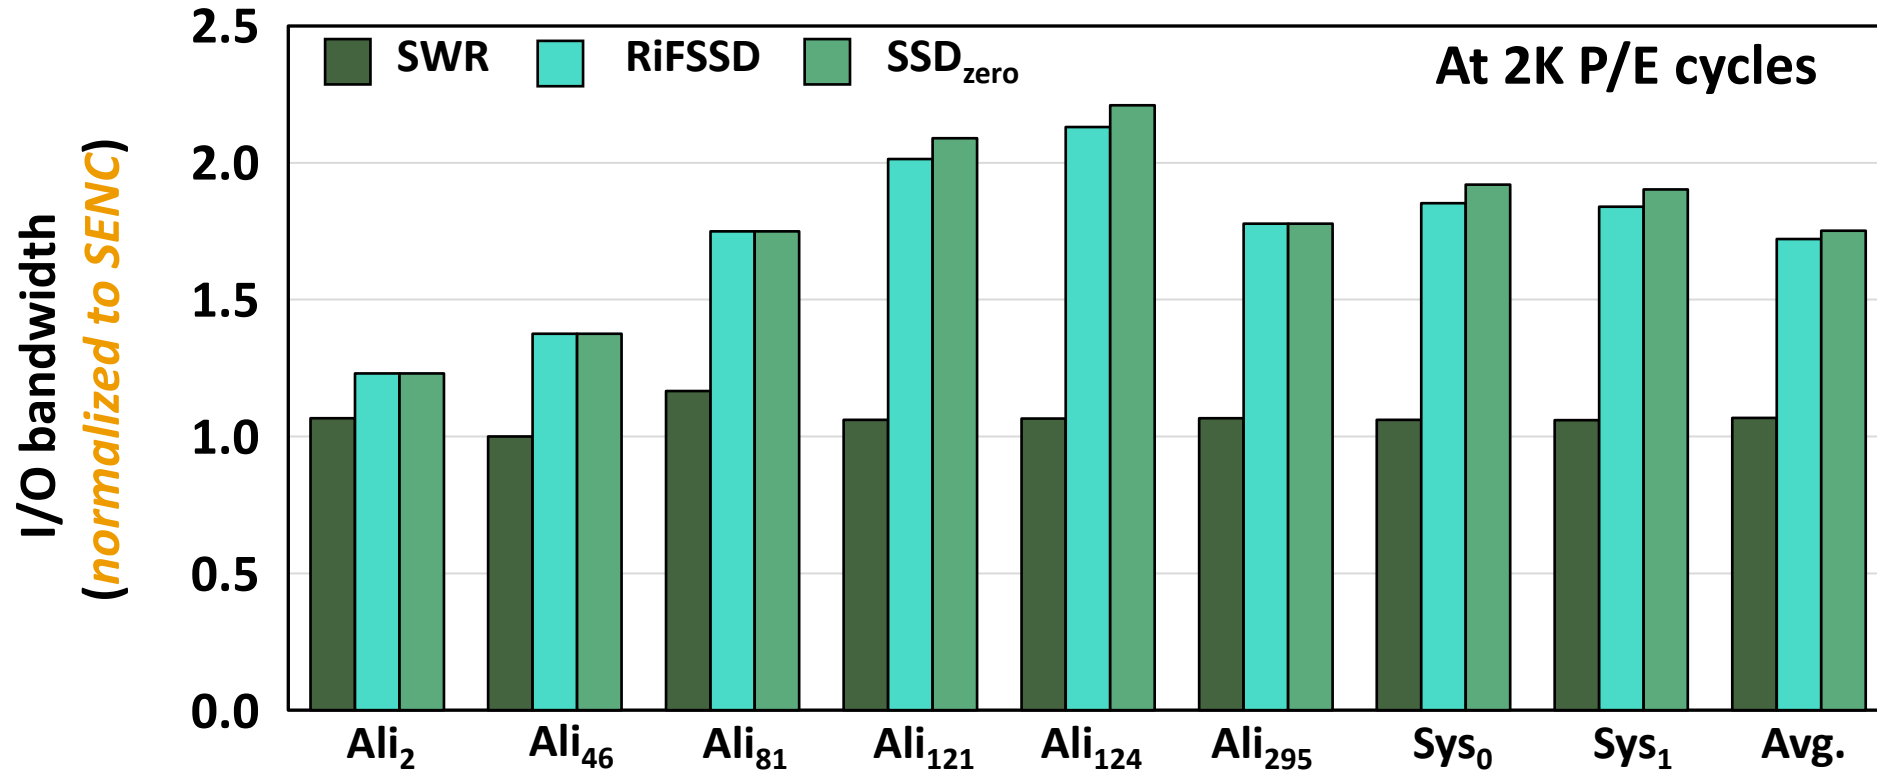
004

Evaluation Results and Conclusion

# Experiment Settings

- ***Simulator:*** MQSim-E
  - Extend NAND flash models w/ real-device characterization results
- ***Workloads:*** 8 real-world traces
  - 6 from AliCloud traces
  - 2 from Systor traces
- ***Comparison***
  - **SSD<sub>zero</sub>**: An ideal SSD where no read-retry occurs
  - **SENC**: A state-of-the-art read-retry mitigation scheme (Sentinel, MICRO21)
  - **SWR**: A state-of-the-art read-retry mitigation scheme (Swift-Read, ISSCC22)

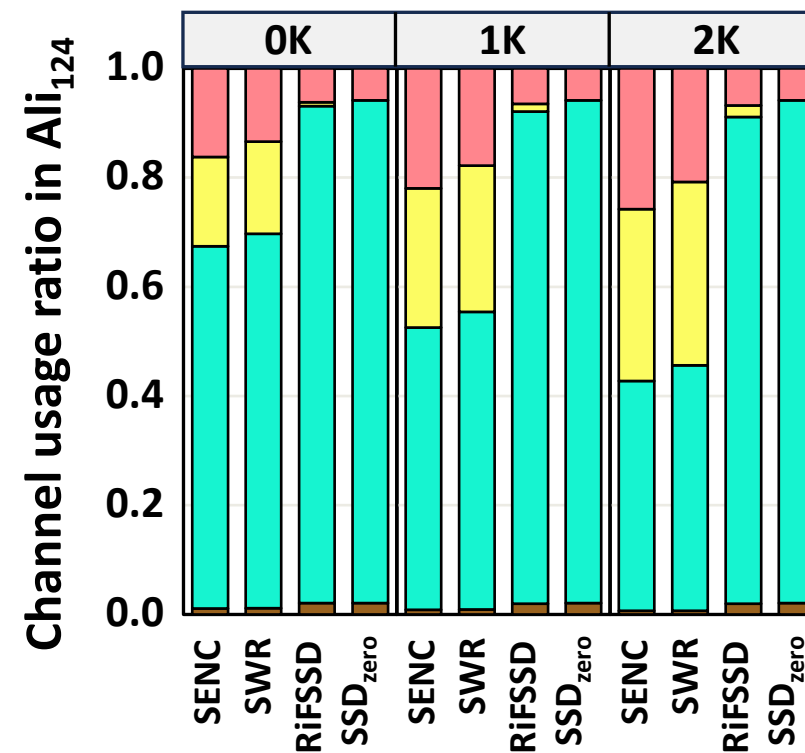
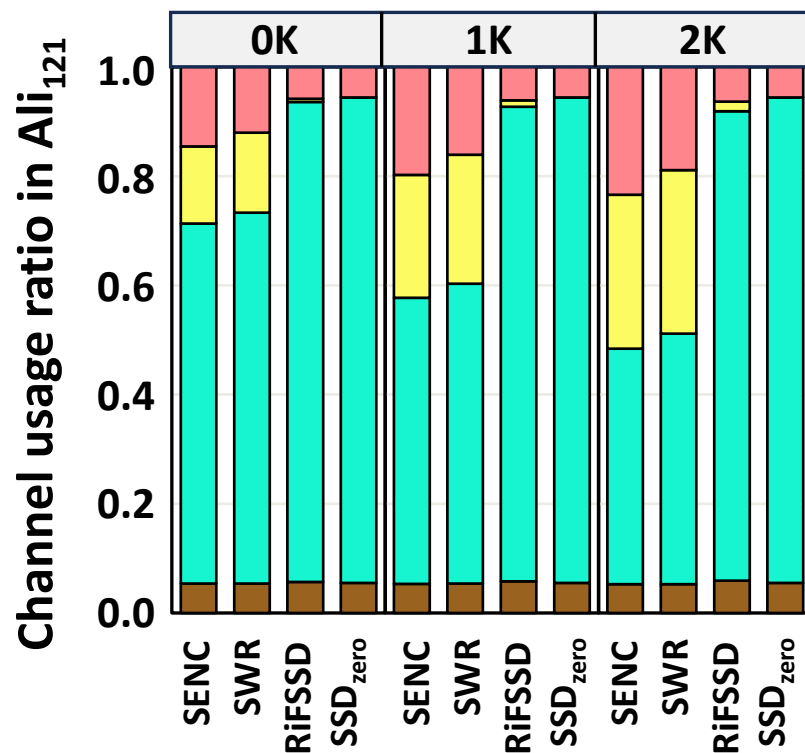
# Experiment Results



RiF improves the I/O bandwidth **71.2%, 61.2%**  
over SENC, SWR

# Experiment Results (cont'd)

■ IDLE: the time when flash channel is no in use  
■ COR: the time to transferring correctable pages  
■ UNCOR: the time to transferring uncorrectable pages  
■ ECCWAIT: the time spent waiting for previous ECC decoding



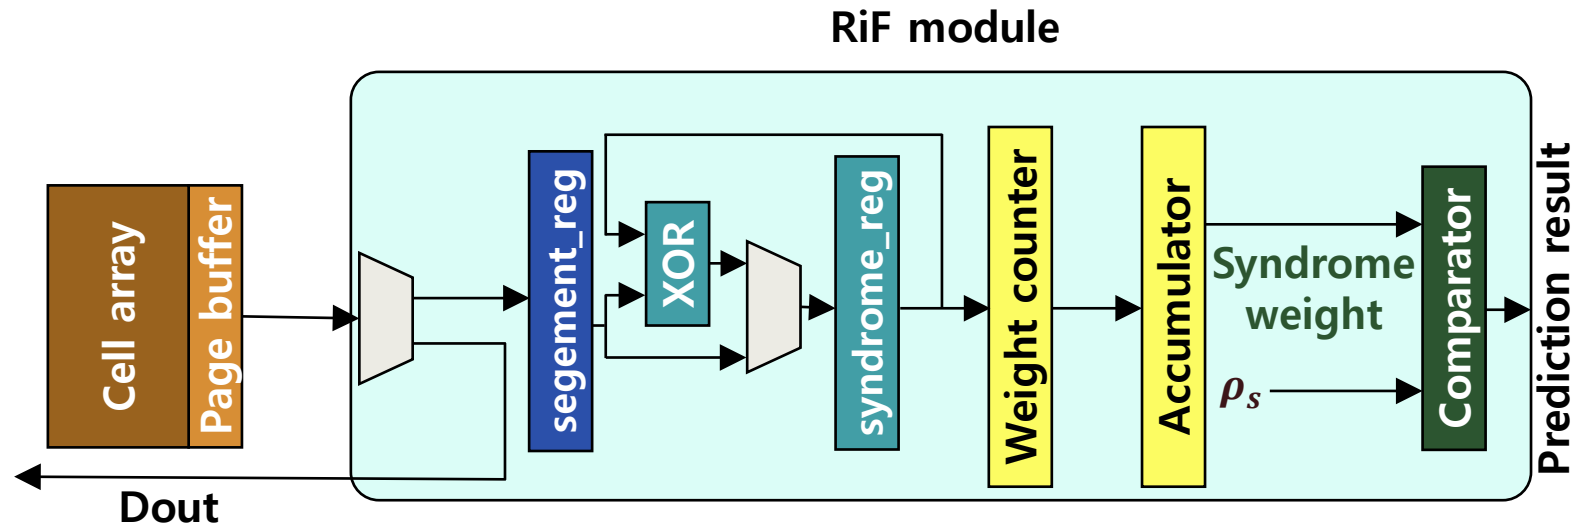
RiF efficiently reduces *wasted channel bandwidth* from *UNCOR, ECCWAIT*

# Conclusion

- Observed that SOTA RR optimizations cannot prevent *the degradation of effective channel bandwidth*
- Presented the RiF scheme, which *determines early on whether a read-retry is required at the flash-chip level*
- Presented a highly optimized RiF module that uses *syndrome weight-based RR prediction*
- Showed the RiF scheme *improved SSD bandwidth by 72.1% on average at 2K P/E cycles*

**Thanks for Listening!**  
**Any Questions? 😊**

# Overhead Analysis



- ✓  **$2.5 \mu s$**  time overhead
- ✓  **$0.012 mm^2$**  area consumption
- ✓  **$1.28 mW$**  power consumption

# Need for Randomization

